

# Introduction to Android



**Mariusz Nowostawski**  
Department of Information Science  
University of Otago, Dunedin, New Zealand

<http://www.business.otago.ac.nz/infosci/>

# What is Android?

- A full software stack
  - Operating system
  - Middleware
  - Application framework
  - Key applications and services



# Android: features

- Rich development environment: device emulator, tools for debugging, memory and performance profiling, plugin for the Eclipse IDE
- Application framework: enabling reuse and replacement of components
- Dalvik Virtual Machine
- Browser, optimized graphics: custom 2D graphics library; 3D graphics based on the OpenGL ES 1.0 specification, SQLite for structured data storage, media support

# Features

- Rich development environment: device emulator, tools for debugging, memory and performance profiling, plugin for the Eclipse IDE
- Application framework: enabling reuse and replacement of components
- Dalvik Virtual Machine
- Browser, optimized graphics: custom 2D graphics library; 3D graphics based on the OpenGL ES 1.0 specification (hardware acceleration optional), SQLite for structured data storage, media support

## APPLICATIONS

Home

Contacts

Phone

Browser

...

## APPLICATION FRAMEWORK

Activity Manager

Window Manager

Content Providers

View System

Package Manager

Telephony Manager

Resource Manager

Location Manager

Notification Manager

## LIBRARIES

Surface Manager

Media Framework

SQLite

OpenGL | ES

FreeType

WebKit

SGL

SSL

libc

## ANDROID RUNTIME

Core Libraries

Dalvik Virtual Machine

## LINUX KERNEL

Display Driver

Camera Driver

Flash Memory Driver

Binder (IPC) Driver

Keypad Driver

WiFi Driver

Audio Drivers

Power Management

# The bottom layer

---

- Linux version 2.6 for core system services such as security, memory management, process management, network stack, and driver model.
- The kernel acts as an abstraction layer between the hardware and the rest of the software stack.
- New open LIBC core

# Application framework

- Full access to the same framework APIs used by the core applications.
- The application architecture is designed to simplify the reuse of components; any application can publish its capabilities and any other application may then make use of those capabilities.
- This mechanism allows components to be replaced by the developer/user.

# Framework: core components

---

- no main() function
  - stress on component reuse
- 
- Activities
  - Services
  - Broadcast receivers
  - Content providers

# Activities

- Visual representations for user interactions
- Consists of a single or multiple Views
- A single application can have single or multiple activities
- Activities are typically independent

# Services

---

- Represent tasks run on behalf of a user that do not have visual interface
- E.g. media player
- It is possible to connect to ongoing service
- Or start service when needed

# Broadcast receivers

---

- Consume and react to system and application-level broadcasts
- E.g. battery low indication
- They do not have visual representation, but can start Activity in response to the broadcast
- Notification manager is a system-level handler for broadcasts

# Content providers

---

- Make a specific set of the application's data available to other applications
- SQLite or filesystem can be used to store persistent data
- Triggered through ContentResolver

# Intents

- Activities, Services and Broadcast receivers are triggered through asynchronous messages:  
*INTENTS*
- For activities and services intents contain action name being requested and specifies the URI of the data to act on.
- For broadcast receivers, the Intent object names the action being announced. For example, it might announce to interested parties that the camera button has been pressed.

# Activities and Tasks

---

- Task is a stack of activities. A sequence of consecutive activities run by the user.
- The activities may belong to the same or separate applications – for the user it will always appear as if the task forms a uniform application.
- E.g. navigation application.

# Framework services

- Content Providers – enable applications to access data from other applications (such as Contacts), or to share their own data
- Resource Manager – provides access to non-code resources such as localized strings, graphics, and layout files
- Notification Manager – enables all applications to display custom alerts in the status bar
- Activity Manager – manages the lifecycle of applications and provides a common navigation backstack

# Library support

- Libc subset, BSD-based, tuned for embedded use
- No Swing, no AWT, no standard Java packages
- Android-specific equivalents for everything
- Media Libraries, many media codecs
- Surface Manager - manages access to the display subsystem and seamlessly composites 2D and 3D graphic layers from multiple applications
- 3D libraries - an implementation based on OpenGL ES 1.0 APIs; FreeType - bitmap and vector font rendering
- SQLite

# Programming mobile devices

---

- Priority: battery life.
- Remember: anytime instruction is executed, it consumes power.
- Processing time scales
  - Human interaction scale: 10-30 interactions/sec
  - Human perception scale: 25-30 frames/sec
  - Computer scale: as much and as fast as possible

# Rules of thumb

---

- Avoid computer scale computations (!)
- Spend most of its time sleeping
- React quickly and decisively to user and network input
- Avoid object allocations
- Short lived objects should be avoided because of GC costs. However - long lived objects take up memory. Trade-offs.

# Android – Dalvik VM

- Work on Android has started in 2005
- Dan Bornstein – lead developer
  - experienced VM architect <http://www.milk.com/home/danfuzz/>
- 6 core developers plus contributors (libs)
- Dan developed the core of Dalvik VM
- Apache Harmony (subset) used for core libs

# Dalvik – the name – the village



- Iceland's north-most fishing village
- 2000 inhabitants
- close to Arctic Circle

# Dalvik – the name

---



# Design objectives

---

- Slow CPU (initial target ARM architecture)
- Limited RAM
- OS does not have any swap space
- Takes battery life into consideration

# CPU considerations

---

- Speed: 250-500 MHz
  - e.g. Qualcomm MSM 7200A 528 MHz
- Bus speed: 100MHz
- Data cache: 16-32kB

# Memory considerations

- Memory efficiency top priority
- Design
  - total system RAM – 64 MB
  - Available RAM after low-level startup: 40MB
  - Available RAM after high-level services started: 20MB
- Multiple independent processes, separate address spaces, separate memory
- Added difficulty: large system library (10MB)

# Bytecode optimisations

---

- Objectives:
  - Single constant pool, minimal repetitions
  - Per-type pools (implicit typing)
  - Implicit labeling

dx tool:

Java bytecode (.class) → Dalvik Executable (.dex)

# Bytecode optimisations

---

- Results
  - on average DEX files are 50% smaller than uncompressed JAR files
  - on average DEX files are just slightly smaller than compressed JARs
- Why all the fuss then?

# Other optimisations

---

- Byte swapping, padding (unnecessary on ARM)
- Static linking
- Inlining special native methods
- Pruning empty methods, e.g. constructor  
Object()
- Adding auxiliary data

# Bytecode verification

---

- DEX structures cannot lie, and
  - Code can't misbehave
  - DEX structures: valid indices, valid offsets
- 
- The verification is done mostly for the sake of bugs prevention (icing on the cake)
  - Security is handled by ... Linux.

# No Just-in-Time (JIT) compiler

---

- Java and Dalvik designed mostly for user apps
- JIT wouldn't work well (warm-up time)
- JIT would put extra pressure on CPU/Mem
- Native code for libraries, graphics, media
- Hardware support: graphics, audio, etc
- For performance conscious there is JNI support

# Dalvik memory model

---

- Clean memory (read only)
  - common DEX files (libraries and sys. libraries)
  - application-specific DEX files
- Shared dirty (mostly not written)
  - library ‘live’ DEX structures
  - shared copy-on-write heap
- Private dirty (read-write)
  - application heap, app. live DEX structures

# GC design

- Considerations
  - Separate processes, heaps, and separate GCs
  - GCs must be independent
  - GC should respect the sharing
- Two generic strategies
  - GC data is stored together with the objects
  - GC data is stored separate from the objects

# GC design

- Keeping GC data separate
  - Avoids un-sharing pages
  - Exhibits better small cache behaviour
  - Does not waste memory

# Instruction set

---

- Java? Not quite. Not at all!
- Completely new instruction set
- Register machine (infinite register machine)
  - Avoids unnecessary memory accesses (tell how?)
  - Consumes instruction stream more efficiently
  - Higher semantic density per instruction (!)
- Result: 30% fewer instructions, 35% fewer code units, 35% more bytes in the instruction stream

# Example

```
public static long sumArray(int[] arr) {  
    long sum = 0;  
    for (int i : arr) { sum += i; }  
    return sum;  
}
```

	bytes	dispatches	reads	writes
.class	25	14	45	16
.dex	18	6	19	6

# Instruction dispatch

```
static void interp(const char* s) {  
    for (;;) {  
        switch (* (s++)) {  
            case 'a': printf ("Hello"); break;  
            case 'b': printf (" "); break;  
            case 'c': printf ("world!"); break;  
            case 'd': printf ("\n"); break;  
            case 'e': return;  
        }  
    }  
}  
int main (int argc, char** argv) { interp("abcde"); }
```

# Computed GOTO (GCC way)

```
#define DISPATCH() { goto *op_table[*((s)++) - 'a']; }

static void interp(const char* s) {
    static void* op_table[] =
        { &&op_a, &&op_b, &&op_c, &&op_d, &&op_e };
    DISPATCH();
    op_a: printf("Hello"); DISPATCH();
    op_b: printf (" "); DISPATCH();
    op_c: printf ("world!"); DISPATCH();
    op_d: printf ("\n"); DISPATCH();
    op_e: return; }
```

# Best dispatch implementation

---

- The computed GOTO can be further optimised if we re-write it in assembly.
- The code above uses typically two memory reads. We can lay out all our bytecodes in memory in such a way that each bytecode takes exactly the same amount of memory - this way we can calculate the address directly from the index.
- Added benefit is the cacheline warm-up for frequently used bytecodes.

# References

---

- Dalvik on Wikipedia  
[http://en.wikipedia.org/wiki/Dalvik\\_virtual\\_machine](http://en.wikipedia.org/wiki/Dalvik_virtual_machine)
- Android resources  
<http://nzdis.otago.ac.nz/projects/wiki/android>
- Google I/O 2008 Talk (Dan Bornstein)  
<http://www.youtube.com/watch?v=ptjedOZEXPM>

# Thank you



mariusz@nowostawski.org  
twitter: praeteritio  
skype: nowostawski

<http://www.otago.ac.nz/>